



DALHOUSIE
UNIVERSITY

CSCI 5411 – Mid Term Project

Quiz UP

Study Notes & Quiz Generator with RAG on AWS

Name: Shray Moza

ID: B00987463

1) Domain Selection & Justification

App Name: Quiz up

Domain: Academic productivity with the help of **note generation** and **quiz creation** from lecture PDFs or notes.

Problem Statement: Students spend most of their time summarizing lecture material and creating practice questions to get a better understanding. Manual effort is inefficient and prone to missing key concepts. We want to create a system that ingests PDFs/notes, extracts core ideas, generates concise study notes, and produces formative assessments.

Business Case: - Increases student productivity and retention by converting raw materials into structured notes and quizzes. Supports retrieval of relevant content on demand with a Retrieval-Augmented Generation (RAG) workflow.

For the implementation I will be using Amplify, Cognito, SageMaker, S3, Lambda, and API Gateway and Cloudwatch. These services are sufficient to support hosting, inference, and secure access.small open-source models and a pragmatic vector search, enabling a fully in-house pipeline.

2) AI Approach & Techniques

1. RAG pipeline
 2. Embeddings with MiniLM
 3. FAISS vector store
-

3) Functional Requirements


1. **Upload & Storage:**
 - ❖ Users upload PDFs.
 - ❖ System stores originals and outputs in S3.
2. **Preprocessing:** Extract, clean, and chunk text.
3. **Embedding & Indexing:**
 - ❖ Compute embeddings for chunks.
 - ❖ Build FAISS index.
 - ❖ Persist to S3.
4. **Query / Task Types:**
 - ❖ Study notes for a document/topic.
 - ❖ A 5-question quiz (MCQ) with an answer key.
5. **RAG Pipeline:**
 - ❖ Retrieve top-k relevant chunks.
 - ❖ Compose prompt with context.

- ❖ Call model for generation.
 - 6. **Delivery:** Return generated notes/quiz to user and store artifacts in S3.
 - 7. **Monitoring:** Basic metrics and logs via CloudWatch (Lambda + SageMaker endpoint/notebook).
 - 8. **Auth:** Cognito user pool for authenticated access.
-

4) Non-Functional Requirements

- **Security:**
 - ❖ IAM least privilege.
 - ❖ Private S3 buckets.
 - ❖ VPC for SageMaker endpoint.
 - ❖ HTTPS via API Gateway.
 - ❖ Cognito for authenticated calls.
 - **Scalability:**
 - ❖ Serverless API (Lambda/API Gateway).
 - ❖ SageMaker endpoint scales by instance size.
 - ❖ Notebook-only dev operations for batch jobs.
 - **Reliability:**
 - ❖ S3 durability for inputs/outputs.
 - ❖ Index stored redundantly in S3.
 - ❖ Retriable Lambda steps.
 - **Performance:**
 - ❖ End-to-end generation target < 5–8s for small inputs.
 - ❖ Batch builds for large PDFs.
 - ❖ Caching common prompts.
 - **Cost Efficiency:**
 - ❖ Small instance classes (ml.t3.medium, ml.c5.large).
 - ❖ Endpoint up only during demos.
 - ❖ Notebooks stopped when idle.
 - ❖ CloudWatch log retention controls.
 - **Observability:**
 - ❖ CloudWatch metrics/alarms for Lambda errors and SageMaker invocation latency.
 - ❖ S3 access logs enabled.
-

5) Initial Architecture Design (made with draw.io)




5.1 Services

1. **S3** – durable storage for PDFs, indexes, and generated outputs.
2. **Amplify** – Hosting the web app, easy auth integration, CI/CD
3. **SageMaker Notebook** – orchestrates data prep, embedding, and FAISS index build.
4. **FAISS (in Notebook/Code)** – vector search (no managed vector DB required).
5. **SageMaker Realtime Endpoint** – hosts a small open-source text-generation model for consistent inference via HTTPS.
6. **Lambda** – RAG API: loads FAISS index from S3, retrieves chunks, builds prompts, calls SageMaker endpoint.
7. **API Gateway** – HTTPS interface for UI/clients.
8. **Cognito** – authentication and JWT validation at API Gateway.
9. **CloudWatch** – logs/metrics/alarms.

5.2 Architecture Description

1. User opens webapp hosted on **Amplify**.
 2. User authenticates using **Cognito**.
 3. User uploads PDF via UI → **API Gateway** → **Lambda** → stores in **S3**.
 4. Processing job triggered: **SageMaker Notebook** reads PDF from **S3**, extracts & chunks text, computes embeddings, builds **FAISS** index, and saves index.faiss + docstore.pkl to **S3**.
 5. On generation request: **Lambda** loads FAISS index from **S3**, retrieves top-k chunks for the prompt, calls **SageMaker Endpoint** with prompt+context, returns generated notes/quiz to UI, and stores outputs in **S3**.
 6. **CloudWatch** collects logs/metrics.
-

6) Data Sequence Diagram (made with draw.io)



7) AWS Services & Tech Stack

1. **Compute/ML:** SageMaker Notebook, SageMaker Realtime Endpoint (Flan-T5-Base default)
 2. **Vector Search:** FAISS (Python), artifacts in S3
 3. **Embeddings:** sentence-transformers all-MiniLM-L6-v2
 4. **API:** API Gateway, Lambda
 5. **UI:** Amplify
 6. **Storage:** S3 with separate prefixes for raw, index, outputs
 7. **Auth:** Cognito user pool & authorizer
 8. **Monitoring:** CloudWatch Logs, basic alarms on error rates/latency.
-

8) Security Analysis

1. **IAM Least-Privilege:** Separate roles for Notebook, Endpoint, Lambda. Each has scoped permission:
 - ❖ Lambda can GetObject from indexes/ and raw/ prefixes and PutObject to outputs/ only.
 - ❖ Notebooks can read raw/, write indexes/.
 - ❖ Endpoint can read models from ECR/S3 as needed.
 2. **S3 Hardening:**
 - ❖ Block public access.
 - ❖ Bucket policies restricted specific IAM roles.
 - ❖ Server-side encryption (SSE-S3).
 - ❖ Enable access logs.
 3. **Network:** Security groups to limit access.
 4. **API:**
 - ❖ HTTPS via API Gateway.
 - ❖ Cognito authorizer.
 - ❖ Request/response size limits.
 5. **Secrets:**
 - ❖ No hard-coded keys.
 - ❖ Rely on IAM roles.
 6. **Data Handling:**
 - ❖ PDFs and outputs stored under per-user prefixes.
 - ❖ Redact PII in logs.
-

9) Identification of Potential Architectural Challenges

1. **Instance Limits:**
 - ❖ Learner Lab restricts instance families/sizes.
 - ❖ Large models won't deploy → use small instruction-tuned models.

2. **Service Availability:** OpenSearch/others may be unavailable → FAISS-on-S3 fallback is primary.
 3. **Endpoint Cost & Warmup:**
 - ❖ Keep endpoints off when not in use.
 - ❖ Pre-warm before demos.
 4. **Latency:** Loading FAISS from S3 per request may add latency → cache index in memory via provisioned Lambda or warm start pattern.
 5. **Token Limits:** Use chunking and concise prompts to return summaries section-wise for long inputs.
-

10) Initial Cost Estimation (conservative, 1–2 weeks of development/use)

Service	Assumption	Est. Cost
SageMaker Notebook	ml.t3.medium, few hours/day, stopped when idle	\$3–\$8
SageMaker Endpoint	ml.t3.medium, used 4–5 hrs/day for 5–7 days	< \$15
Lambda + API Gateway	Low traffic dev/test	<\$3
S3 (data in GBs) + Transfer	PDFs + indexes + outputs	<\$2
CloudWatch Logs	Basic retention, low volume	\$1–\$2
Total		< 40
